



## PROJECT

## Generate TV Scripts

A part of the Deep Learning Nanodegree Foundation Program

## PROJECT REVIEW

## CODE REVIEW

## NOTES

SHARE YOUR ACCOMPLISHMENT!  

## Requires Changes

3 SPECIFICATIONS REQUIRE CHANGES

## Required Files and Tests

The project submission contains the project notebook, called "dlnl\_tv\_script\_generation.ipynb".

All the unit tests in project have passed.

Great job! All tests have passed 🍌

## Preprocessing

The function `create_lookup_tables` create two dictionaries:

- Dictionary to go from the words to an id, we'll call `vocab_to_int`
- Dictionary to go from the id to word, we'll call `int_to_vocab`

The function `create_lookup_tables` return these dictionaries in the a tuple (`vocab_to_int`, `int_to_vocab`)

Nice job, but it is more Pythonic to construct these using dictionary comprehensions:

```
vocab_to_int = {word: ii for ii, word in enumerate(set(text))}
int_to_vocab = {ii: word for ii, word in enumerate(set(text))}
```

The function `token_lookup` returns a dict that can correctly tokenizes the provided symbols.

## Build the Neural Network

Implemented the `get_inputs` function to create TF Placeholders for the Neural Network with the following placeholders:

- Input text placeholder named "input" using the TF Placeholder name parameter.
- Targets placeholder
- Learning Rate placeholder

The `get_inputs` function return the placeholders in the following the tuple (Input, Targets, LearningRate)

The `get_init_cell` function does the following:

- Stacks one or more BasicLSTMCells in a MultiRNNCell using the RNN size `rnn_size`.
- Initializes Cell State using the MultiRNNCell's `zero_state` function
- The name "initial\_state" is applied to the initial state.
- The `get_init_cell` function return the cell and initial state in the following tuple (Cell, InitialState)

Nice job adding the dropout wrapper. Additional LSTM layers here will allow the capturing of more complex sequence representations.

The function `get_embed` applies embedding to `input_data` and returns embedded sequence.

Nice work; you can also use `tf.contrib.layers.embed_sequence(input_data, vocab_size, embed_dim)` which maps a sequence of symbols to a sequence of embeddings.

The function `build_rnn` does the following:

- Builds the RNN using the `tf.nn.dynamic_rnn`.
- Applies the name "final\_state" to the final state.
- Returns the outputs and final\_state state in the following tuple (Outputs, FinalState)

The `build_nn` function does the following in order:

- Apply embedding to `input_data` using `get_embed` function.
- Build RNN using cell using `build_rnn` function.
- Apply a fully connected layer with a linear activation and `vocab_size` as the number of outputs.
- Return the logits and final state in the following tuple (Logits, FinalState)

Good. You can also include weight and bias initializers here.

The `get_batches` function create batches of input and targets using `int_text`. The batches should be a Numpy array of tuples. Each tuple is (batch of input, batch of target).

- The first element in the tuple is a single batch of input with the shape [batch size, sequence length]
- The second element in the tuple is a single batch of targets with the shape [batch size, sequence length]

Well done

## Neural Network Training

- Enough epochs to get near a minimum in the training loss, no real upper limit on this. Just need to make sure the training loss is low and not improving much with more training.
- Batch size is large enough to train efficiently, but small enough to fit the data in memory. No real "best" value here, depends on GPU memory usually.
- Size of the RNN cells (number of units in the hidden layers) is large enough to fit the data well. Again, no real "best" value.
- The sequence length (`seq_length`) here should be about the size of the length of sentences you want to generate. Should match the structure of the data. The learning rate shouldn't be too large because the training algorithm won't converge. But needs to be large enough that training doesn't take forever. Set `show_every_n_batches` to the number of batches the neural network should print progress.

You should use enough epochs to get near a minimum in the training loss. Use enough epochs so you get low training loss, where the loss isn't improving much with additional training.

Your RNN size is too small. To be able to generalize better, you'll likely need to try values over 100. Try 64, 128, 256, and 512 and see what works best. Think, from the top of the notebook, the dataset has almost 12k unique words, over 200 different scenes, and over 4000 lines. You are using a single hidden layer with only 16 nodes. Intuitively, do you think 16 nodes will be able to capture a lot of information present in 12000 words?

Similarly, your embedding dimension is small. There are a few approaches to selecting this. I would try a wide range of values, say 16, 256, 512, and 1024 and see which one performs best. I recommend reviewing the lessons on word embeddings to get a stronger intuition about the purpose of this step.

Lastly, you may want to slightly decrease your learning rate if you find that your training loss won't decrease.

The project gets a loss less than 1.0

You need a loss less than 1.0.

```
train_loss = 5.248
```

## Generate TV Script

"input:0", "initial\_state:0", "final\_state:0", and "probs:0" are all returned by `get_tensor_by_name`, in that order, and in a tuple

The `pick_word` function predicts the next word correctly.

Good. This will ensure that your network doesn't fall into prediction loops.

The generated script looks similar to the TV script in the dataset.

It doesn't have to be grammatically correct or make sense.

This does not look like the underlying script. For example, it doesn't learn the relationship between a new line and someone speaking. This is evident in this line:

```
ould show(homer and crossed stupid evil joey_kramer: bag for wood talk-sings, i buy go to can
```

Additionally, there is a strange use of punctuation persistent throughout the prediction.

 RESUBMIT

 [DOWNLOAD PROJECT](#)



## Best practices for your project resubmission

Ben shares 5 helpful tips to get you through revising and resubmitting your project.

 [Watch Video](#) (3:01)

[RETURN TO PATH](#)